

# Lab 1 Report

Alp Can Aloglu

SID : 862190749

E-mail: aalog001@ucr.edu

- **All code changes you made on xv6 source code** (hint: `diff -r original_xv6 your_xv6`)
- **Detailed explanation about these changes**

Changes in sysproc.c:

```
[aalog001@sledge ~]$ diff -r xv6_original/sysproc.c xv6/sysproc.c
19,20c19,25
< exit();
< return 0; // not reached
---
> // Copied Kill sys_kill
> int status;
passing the user-level argument to the kernel-level argument we use argint()
> if(argint(0, &status)< 0){
>     return -1;
> }
> exit(status); used the new signature of exit function
> return 0; // not reached
26c31,38
< return wait();
---
> // Found a description on how to pass struct pointer to system call
> // https://stackoverflow.com/questions/53383938/pass-struct-to-xv6-system-call
> // used this example.
> int *status; declared the status pointer of type int
Passing arguments from user-mode to kernel-mode using argint()
> if(argptr(0,(void*)&status,sizeof(*status))<0){>     return -1;
> }
> return wait(status);used the new signature of wait function
91a104,126
>
> // Wait pid implementation
> int
> sys_waitpid(void)
> {
>     int pidToWait,options; declared variables for the process pid and option of type int
>     int *status ;declared the status pointer of type int
```

## Lab 1 Report

Passing arguments from user-mode to kernel-mode by using argint() for type int and argptr() for type of pointer

```
> if(argint(0,&pidToWait)< 0)
>     return -1;
> if(argptr(1,(void *)&status,sizeof(*status)) < 0)
>     return -1;
> if(argint(2,&options)< 0)
>     return -1;
> return waitpid(pidToWait, status, options);Called the waitpid function in Kernel mode
> }
```

Changes in proc.c:

```
[aalog001@sledge ~]$ diff -r xv6_original/proc.c xv6/proc.c
228c228
< exit(void)
---
> exit(int status) We change the function signature by adding the argument status of type int
233a234,236
> // Set exit status
> curproc->exitStatus = status; Adding this line of code in order to set the current process's status to the status passing in to the exit function
>
273c276
< wait(void)
---
> wait(int *status) Changing the function signature by adding the argument status of type int pointer
278c281
<
---
>
288a292,293
If the process goes to the ZOMBIE state:
>     if(status) Checking if the status is not nullptr
>         *status = p->exitStatus; If the status pointer is not null, then we change it to the process's status
534a540,591
>
```

## Lab 1 Report

```
> // Waits for specific process to finish
> int
> waitpid(int pidToWait, int *status,int options)
> {
>     struct proc *p;
>     int haveproc, pid; Changing havechild to haveproc because in waitpid function we do not care about the child process but we care about an specific process with a given pid.
>     struct proc *curproc = myproc();
>
>     acquire(&ptable.lock);
>     for(;;){
>         // Scan through table looking for exited process with given pid
>         haveproc = 0;Changing havechild to haveproc
>         for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
>             if(p->pid != pidToWait)
>                 continue;
>             haveproc = 1;Changing havechild to haveproc
>             if(p->state == ZOMBIE){
>                 // Found one.
>                 If the process goes to the ZOMBIE state:
>                 // Status null check
>                 if(status)Checking if the status is not nullptr
>                     *status = p->exitStatus; If the status pointer is not null, then we change it to the process's status
>                     pid = p->pid;
>                     kfree(p->kstack);
>                     p->kstack = 0;
>                     freevm(p->pgdir);
>                     p->pid = 0;
>                     p->parent = 0;
>                     p->name[0] = 0;
>                     p->killed = 0;
>                     p->state = UNUSED;
>                     release(&ptable.lock);
>                     return pid;
>                 }
>             }
>
>             // No point waiting if we don't have any process with given pid.
>             // Also if current process is killed already, no point waiting.
>             if(!haveproc || curproc->killed){Changing havechild to haveproc
>                 release(&ptable.lock);
>                 return -1;
>             }
>         }
>     }
> }
```

## Lab 1 Report

```
>      }
>
>      // Wait for process to exit. (See wakeup1 call in proc_exit.)
>      sleep(curproc, &phtable.lock); //DOC: wait-sleep
>  }
> }
```

### Changes in defs.h:

```
[aalog001@sledge ~]$ diff -r xv6_original/defs.h xv6/defs.h
107c107
< void      exit(void);
---
> void      exit(int);Modified exit forward declaration
120c120
< int      wait(void);Modified wait forward declaration
---
> int      wait(int*);
122a123,124
> int      waitpid(int,int*,int); // Added waitpid forward declaration
```

### Changes in user.h:

```
[aalog001@sledge ~]$ diff -r xv6_original/user.h xv6/user.h
6,7c6,7
< int exit(void) __attribute__((noreturn));
< int wait(void);
---
> int exit(int) __attribute__((noreturn)); // Modified exit
> int wait(int*); //Modified wait
25a26,27
> int waitpid(int,int*,int); // Added waitpid
```

### Changes in usys.S:

```
diff -r xv6_original/usys.S xv6/usys.S
31a32,33
> SYSCALL(waitpid) // Added waitpid to SYSCALL list
```

## Lab 1 Report

```
diff -r xv6_original/wc.c xv6/wc.c
```

Changes in syscall.h:

```
diff -r xv6_original/syscall.c xv6/syscall.c
105a106,107
> extern int sys_hello(void); //Toy
> extern int sys_waitpid(void); // Added waitpid
128a131,132
> [SYS_hello] sys_hello, //Toy
> [SYS_waitpid] sys_waitpid, // Added waitpid

diff -r xv6_original/syscall.h xv6/syscall.h
22a23,24
> #define SYS_hello 22 // toy
> #define SYS_waitpid 23 // Added waitpid
```

Changes in syscall.c:

```
diff -r xv6_original/syscall.c xv6/syscall.c
105a106,107
> extern int sys_hello(void); //Toy
> extern int sys_waitpid(void); // Added waitpid
128a131,132
> [SYS_hello] sys_hello, //Toy
> [SYS_waitpid] sys_waitpid, // Added waitpid
```

- Screenshots about how you run the related program(s) and results

## Lab 1 Report

```
[nmoha034@sledge ~]$ cd xv6
[nmoha034@sledge xv6]$ make qemu-nox
which: no qemu in (/usr/lib64/qt-3.3/bin:/usr/cssshare/bin:/usr/cssshare/sbin:/usr/local/bin:/usr/bin:/usr/local/sbin:/home/csmajs/nmoha034/.local/bin:/home/csmajs/nmoha034/bin)
qemu-system-i386 -nographic -drive file=fs.img,index=1,media=disk,format=raw -drive file=xv6.img,index=0,media=disk,format=raw -smp 2 -m 512
(process:9075): GLib-WARNING **: 10:20:37.197: gmem.c:489: custom memory allocation vtable not supported
?7l
SeaBIOS (version 1.11.0-2.el7)

iPXE (http://ipxe.org) 00:03.0 C980 PCI2.10 PnP PMM+1FF94780+1FED4780 C980

Booting from Hard Disk..xv6...
cpu1: starting 1
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap start 58
init: starting sh
$ test
Child Process 2.
Child Process 2 exit status = 2
Child Process 1.
Child Process 1 exit status = 1
Main Child Process.
Main Child Process exit status = 9
Parent Process.
$ [ ]
```

● Clion 2020.3.4 available  
Update...

```
Booting from Hard Disk..xv6...
cpu1: starting 1
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap start 58
init: starting sh
$ test
Child Process 2.
Child Process 2 exit status = 2
Child Process 1.
Child Process 1 exit status = 1
Main Child Process.
Main Child Process exit status = 9
Parent Process.
$ [ ]
```

test.c:

```
#include "types.h"
#include "user.h"
#include "stat.h"

int main(int argc, char *argv[]){
    int pid = fork();
    int *status = (int *) (1);
    int *status2 = (int *) (1);
    int *status3 = (int *) (1);
    if(pid == 0){
        int pid_c1 = fork();
        if(pid_c1 == 0) {
```

## Lab 1 Report

```
int pid_c2 = fork();

if (pid_c2 == 0){
    printf(1,"Child Process 2.\n");
    exit(2);
}
else{
    waitpid(pid_c2,status3,0);
    printf(1,"Child Process 2 exit status = %d\n ", *status3);
    printf(1,"Child Process 1.\n");
    exit(1);
}
}
else{
    waitpid(pid_c1,status2,0);
    printf(1,"Child Process 1 exit status = %d\n ", *status2);
    printf(1,"Main Child Process.\n");
    exit(9);
}
}
else {
    waitpid(pid,status,0);
    printf(1,"Main Child Process exit status = %d\n ", *status);
    printf(1,"Parent Process.\n");
}
exit(0);
}
```

We have called `fork()` system call three times to create three child processes. After each `fork()` system call, we check the pid, and if it is zero (meaning that the process is a child process) or is not a zero (meaning that the process is a parent process), we specify what to do next accordingly.

For example, we made the first child process (main child) by `pid = fork()`, then we checked on `pid` value. If it is not zero (parent process), we call `waitpid` system call (`waitpid(pid, status, 0)`) which means the parent process should wait for the child process with that specific `pid` (which is the main child process). Then we terminated each child with different exit status, and printed out the corresponding exit status when the `waitpid` of each child returns the exit status of the child in question. For example, the main child process terminated with a status value of 9, and we can see that the test program prints number 9 after the `waitpid` that is waiting for the main child process. It means that the parent process waited for the main child process to end (`exit(9)`), and then resumed and continued the execution.